Reinforcement Learning on Route Planning through Google map for Self-driving System

Dung-Yi, Chao
Department of Mechanical Engineering
Purdue University
West Lafayette, Indiana 47906

Abstract

Self-driving system is a popular and important application of artificial intelligence. We implement Double Q-Learning algorithm in our route planning system and introduce a way for the learning agent to interact with Google map. The result is 31% more energy consumption and three times more duration then the route provided by Google map. However, the result can be improved by applying more complicated neural network architecture, higher resolution navigation and unlimited access to Google map database.

Introduction

Background knowledge

Reinforcement learning (RL) is a planning algorithm involved Markov decision process(MDP). We can think of it as mimic of human making decision. We call human as **agent** and encounter a situation such as seeing menu in a restaurant. We call the restaurant as **environment** and the current position and feeling as **state**. In a restaurant, we make a decision and order a meal where the term, order, is **action**. There is chances that the same meal which you have tried multiple times turns out to be ways too spicy than ever. We call the chance as **state transition probability**. It means that even the same action we make after a same state, the result can be different (spicy) from previous (normal flavor). After the meal, we will comment on this meal or restaurant based on the taste or the dinning experience. We say the taste or experience as **reward**. Reward can be positive, meaning that we really enjoy the food or the service. This catering experience would affect the decision we make next time when we are choosing among several restaurants or meals on the menu. This would gradually form or change the routine we make decision and we describe it as **policy**

The whole process can be simplified in Fig.1 and a few words as the following, under an environment, the agent encounter a state and make an action based on the policy. The environment would lead us to the next state and also give us a reward based on the action. We will make another action and receive another reward and enter the next state and so on. We will learn to form and modify the policy based on a series of state-action-reward procedure.
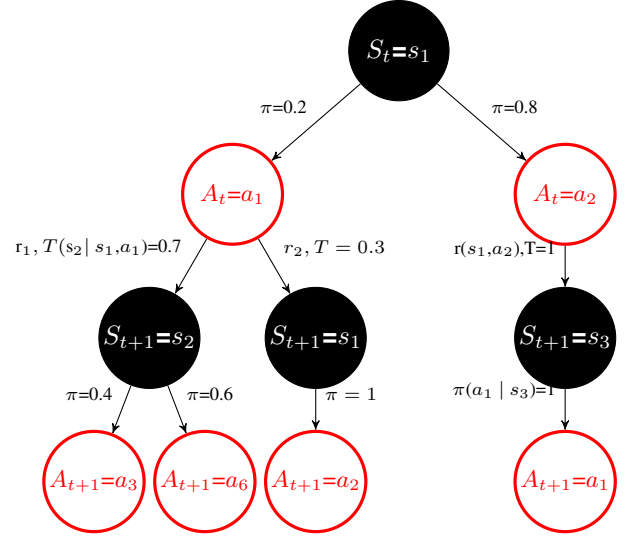

Figure 1: The procedure of iteration between state, action and reward. In state S_t, according to the policy π, we have 0.2 chance to choose action a_1 and 0.8 chance to choose action a_2. The reward r_1 and r_2 can be different.

Some important term and notation we will use in figure 1 and the following paragraph are specified here.

- State set $\mathbf{S} = \{\ s_1, s_2, s_3, ... \}$
- Action set $\mathbf{A} = \{\ a_1, a_2, a_3, ... \}$
- State transition probability function $\mathbf{T} = T(s' \mid$ s,a$) = P[S_{t+1}=s' \mid S_t = s,\ A_t = a]$. This means the probability of transition from state s to s' when taking action a.
- Reward function $\mathbf{r} = $ r(s,a) $= E[R_{t+1} \mid S_t = s,\ A_t = a]$. This means the expected value of reward given **s** and **a**.
- Policy $\pi = \pi($a $\mid$ s$) = P[A_t = a \mid S_t = s]$. This means the probability of choosing **a** given the **s**.
- Discount factor $\gamma \in [0, 1]$

Learning Process

Human learn from the feedback after we make some decision and action. We want the outcome of the ultimate goal as good as possible. For example, a five star restaurant is

located on the top of a mountain. There are bad guys and ferocious animals on our way up to the top. We need to learn a good way to avoid them and get to the top rather than stop at the 2 star restaurant in the middle of the mountain.

Our agent learns from the reward for each action. We assign each action a value and call this action-value function **q(s,a)**. If the agent is at state s and with 4 choices of action, then we get four q value for this state s. The value of q are initially set to 0 and can be updated by the reward which we just encounter, furthermore, we can take the future reward into account in case we might be fooled by the current reward. Here is the functionality of discount factor γ, which weights the importance of the future reward. If γ is near 1, it means the future reward is almost as important as the current reward. We can express the combination of current reward and the future reward by the following equation:

$$q_\pi(s,a) =$$
$$r(s,a) + \gamma \sum_{s' \in S} T(s'|s,a) \sum_{a' \in A} \pi(a'|s') q_\pi(s',a')$$

State s' represents next state relative to current state s. We are not sure which action a' is going to be selected before the agent really in state s', so we need to take expectation of $q_\pi(s',a')$ and that is where the summation and π take the role. Likewise, we are not sure what state s' will the agent enter after taking action a at state s, so we need $T(s'|s,a)$ to represent the possibility and take all the chance into account. After the agent steps in most of the states and tries most of the actions in each state, we can construct an instruction map to demonstrate the quality of taking a specific action given a specific state. At the end, the agent can choose the highest value of action at each state which would most likely to lead us to the optimum result and we refer it to greedy-policy.

We can update $q(s,a)$ through off-policy or on-policy where the update algorithm will be specified in next section. What distinguishes these two policy is how we update $q(s,a)$. For the off-policy, we use the $max q(s',a')$ and $r(s,a)$ to update the $q(s,a)$. For the on-policy, we use $q(s',a')$ and $r(s,a)$ to update the $q(s,a)$.

Value function $v_\pi(s)$ is a function to grade a specific state by taking all the q value in this state into account which shows in the following equation:

$$v_\pi(s) = \sum_{a \in A} \pi(a|s) q_\pi(s,a)$$

Paper Survey

Deep Reinforcement Learning with Double Q-Learning

In this work, the author apply neural network on reinforcement learning and named it Double DQN(Van Hasselt, Guez, and Silver 2016). The functionality of neural network is to map the state s to $q(s,\cdot)$. The input of state s can be an n-dimensional vector such as an image. The output is a m-dimensional vector in which each element can be interpreted as q value corresponds to each action. In short, the neural network is a function to map from $\mathbb{R}^n$ to $\mathbb{R}^m$. We use θ to represent the parameter in the neural network.

They set up a goal or target to evaluate the q value computed by θ. The target is generated by another neural network called target network governed by θ' which contains the same architecture as the neural network governed by θ. Now we get two neural network with identical structure governed by parameters θ and θ' separately. θ' is initially copied from θ. During training in one episode(agent begin from the start to the end), we will copy θ to θ' for every N steps. In other words, we don't update θ' within these N steps. The equation of target is $Y_t \equiv r_{t+1} + \gamma q(S_{t+1}, \underset{a}{argmax} q(S_{t+1}, a; \theta_t), \theta'_t)$. We first input S_{t+1} into the θ network and choose the action a corresponds to the highest value in the output vector denoted by $\underset{a}{argmax} q(S_{t+1}, a; \theta_t)$. At the same time, we input S_{t+1} into the θ' network and get another output vector. We compute the target Y_t by combining the current reward r_{t+1} and the q value from the θ' network $q(S_{t+1}, a; \theta'_t)$. The learning algorithm provided by Zyiu Wang's paper in 2016 is shown in Algorithm 1.

Algorithm 1: Double DQN Algorithm

Input : D-empty replay buffer; θ-initial network parameter; θ'-copy of θ
N_r-replay buffer max size; N_b-training batch size; N-target network update frequency

1 **for** (*episode* $e \in \{$ *1, 2,...,M* $\}$){
2   initialize frame sequence $\mathbf{x} \leftarrow ()$;
3   **for** ($t \in \{$ *0, 1,...* $\}$){
4   Set state $s \leftarrow \mathbf{x}$, sample action $a \sim \pi_B$;
5   Sample next frame x^t from environment ϵ given (s,a) and receive reward r, and append x^t to $\mathbf{x}$;
6   **if** $|\mathbf{x}| > N_f$ **then** delete oldest frame x^{tmin} from $\mathbf{x}$ **end**;
7   Set $s' \leftarrow \mathbf{x}$, and add transition tuples (s, a, r, s') to D,replacing the oldest tuple if $|D| \geq N_r$;
8   Sample a minibatch of N_b tuples $(s, a, r, s') \sim$ Unif(D);
9   Construct target values, one for each of the N_b tuples: Define
$a^{max}(s';\theta) = argmax_{a'} \mathbf{q}(s',a';\theta)$

$$y_j = \begin{cases} r & \text{if } s' \text{ is terminal} \\ r + \gamma \mathbf{q}(s', a^{max}(s';\theta); \theta'), & \text{otherwise.} \end{cases}$$

10   Do gradient descent step with loss $\|y_j - \mathbf{q}(s, a; \theta)\|^2$;
11   Replace target parameters $\theta' \leftarrow \theta$ every N
12   }
13 }

Experience replay is a biological inspired technique to get rid of correlations in the data sequence. We choose data which stored in the replay buffer uniformly at random to compute the loss and updtate the weights during learning.

Value Iteration Network

Value iteration network(VIN) is a model-free planning algorithm(Tamar et al. 2016). We can use VIN with standard backpropagation and RL algorithm to deal with problems required visual perception, continuous control and natural language based decision. The goal is to learn a policy end-to-end which would generalize to solve different, unseen domain.
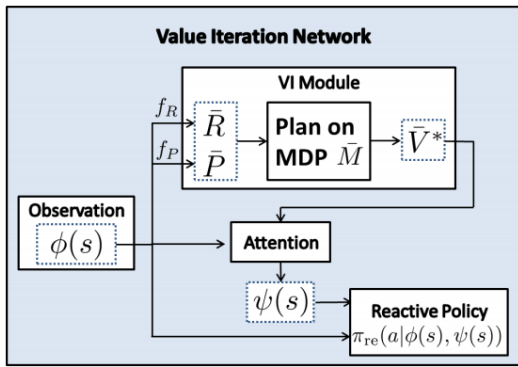
Figure 2: Value iteration network from:Tamar, Aviv, et al. "Value iteration networks." Advances in Neural Information Processing Systems. 2016

Figure 3: Value iteration module from:Tamar, Aviv, et al. "Value iteration networks." Advances in Neural Information Processing Systems. 2016

In figure 2, the input $\phi(s)$ is an image (such as terrain image) and the current state. The output $\pi_{re}(a|\phi(s), \psi(s))$ is a vector of probability over actions. f_R is basically a convolutional neural network (CNN) that transform the input image to a reward image $\bar{R}$(each pixel can represent a reward value). f_P is a state transition function $\bar{P}(s'|s,a)$. $\bar{V}^*$ is a value function which has the same size of $\bar{R}$. Since the optimal policy at state s can depends only on nearby states which are a subset of the $\bar{V}^*$, the author use **Attention** to represent this logic and outputs a vector $\psi(s)$ which represent the value in state set we really care about.

In figure 3, it explains the heart of VIN, VI module, a mechanism to achieve the logic of the equation V(s) = $max_a(R(s,a) + \gamma \sum_{s' \in S} P(s'|s,a)V(s'))$. Each iteration of VI module can be regarded as passing reward image $\bar{R}$ and previous value function into a convolution layer and max-pooling layer and then outputs a new value function $\bar{V}$. Each channel in the convolution layer corresponds to the Q-function for a specific action Q(s,a_1),Q(s,a_2),...Q(s,a_m) where m is the length of action set. The convolution kernel weights correspond to discounted transition probabilities $\gamma P(s'|s,a)$. This layer is then max-pool along the actions channel to produce next iteration of value function $\bar{V}$ where $\bar{V}_{i,j} = max_a \bar{Q}(a,i,j)$. We then attach $\bar{V}$ to $\bar{R}$ as it's second channel and feed them into convolution layer and max-pool layer K times to perform K iterations where K is the minimum value to convey the reward information from the goal to state s. After K iterations, the VI module will output $\bar{V}^*$ for the agent to make decision. We can then apply DQN or other RL methods to train parameters in figure 2 and figure 3.

Generative Image Modeling using Style and Structure Adversarial Networks(S^2-GAN)

Generative Adversarial Networks(GAN) contains two models: generator G and discriminator D. Generator tries to generate images which looks like real image and discriminator tries to distinguish between the real image and generated.

Structure (geometry of scene) and style (texture and illumination) are two key ingredients in image formation while ignored by 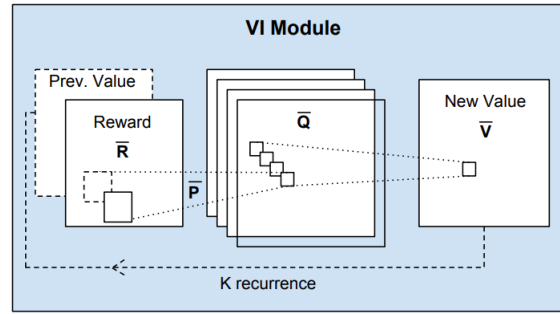most recent generative model. This paper (Wang and Gupta 2016) decompose the generative process into two procedures: (i) use the Structure-GAN to generate a surface normal map and (ii) Style-GAN to take the surface normal map as input and generate the 2D image. The two GANs are trained independently and merged together via join learning.

Structure-GAN The input to the first Generative network is $\hat{z}$ (100-d vector sampled from uniform distribution) and then generate the surface normal map (G($\hat{z}$): 72×72×3). The first Discriminator network takes image(72×72×3) as input and outputs a single value[0,1] which tells the surface normal map is real (closes to 1) or generated (closes to 0).

Style-GAN The input to the second Generative network are $\tilde{z}$ (100-d vector sampled from uniform distribution) and ground truth surface normal and then the network generates images(G(C_i,$\tilde{z}_i$)) with texture and illumination. The input to the second Discriminator network are ground truth surface normal map, G(C_i,$\tilde{z}_i$), real image and real image's surface normal map. This paper also includes fully convolutional network (**FCN**) which takes G(C_i,$\tilde{z}_i$) as input and estimates it's surface normal map in order to make the second Generative network outputs better image aligned with the generated surface normal map.

Joint learning for S^2-GAN After training the Structure-GAN and the Style-GAN independently, we are going to train both networks together shown in figure 4 but first we remove the **FCN** part. Firstly, we input $\hat{z}$ and get generated surface normals G($\hat{z}$) and receive the first loss by feeding G($\hat{z}$) into the Discriminator network in Structure-GAN. Secondly, we input G($\hat{z}$) and $\tilde{z}$ into generator network of Style-GAN and get generated images(G(G($\hat{z}$),$\tilde{z}$)). We now receive the second loss by feeding G(G($\hat{z}$),$\tilde{z}$) and G($\hat{z}$) into the Discriminator network in Style-GAN. We will combine the first loss and the second loss (scaled by 0.1) to train the generator network in Structure-GAN for producing better surface normal map.

We can apply this technique to generate a heat map for the vehicle by input some parameters such as motor output, battery output, ambient temperature and so on. If the heat map is real enough, we can design a better cooling management system and minimize the number of heat sensor. The self-driving system can plan a better route based on the condition of the vehicle.
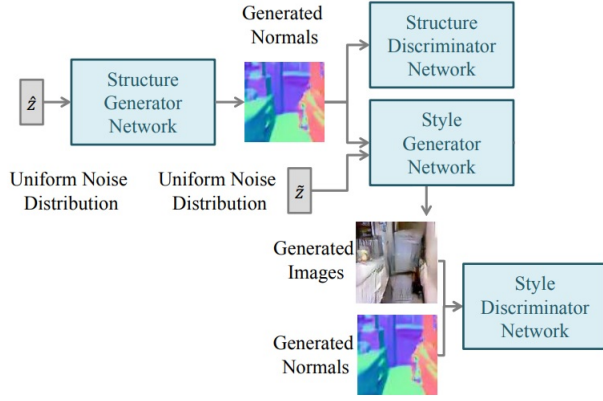
Figure 4: S^2GAN model from:Wang, Xiaolong, and Abhinav Gupta. "Generative image modeling using style and structure adversarial networks." European Conference on Computer Vision. Springer International Publishing, 2016.

Experiments

In this section, we will first introduce the Double-DQN algorithm (Algorithm 1) which we implemented in the experiment through tensorflow library and then explain how do we set up the environment for the learning agent. Lastly, we demonstrate the result of the experiment.

The Learning Agent

Our learning agent is an electric vehicle and navigating on the Google map environment by choosing different action (north, east, south, west). The action can be determined by the Double-DQN or by random. During the learning process, the agent will first navigate on the map randomly to explore the map, but we will gradually reduce the portion of choosing action randomly but adopt the action with highest q value provided by the Double-DQN model. γ is 0.9 and N_b is 32.

The Neural Network Architecture The first layer is an input layer which takes in the geocode of the current position and divided by 180. The input layer is followed by two fully connected layer (10 neurons and 6 neurons respectively) with relu activation function and dropout rate as 0.25. The last layer is a four dimensions output represent the q value for each action. We initialize two identical network, the first one is the Q-network which the agent determines the action at the current state. The second one is the Target-network which acts as a target for Q-network to achieve. We only do backpropagation and update the weights through AdamOptimizer with learning rate 0.0001 in the Q-network for every steps and then copy the weights in Q-network to Target-network for every five steps.

Environment

In order to know if our agent is able to find the best route with minimum energy cost and acceptable duration under the Double DQN algorithm compared to the route provided by Google map API (Geocoding API, Directions API and Elevation API), we set up the experiment as following.
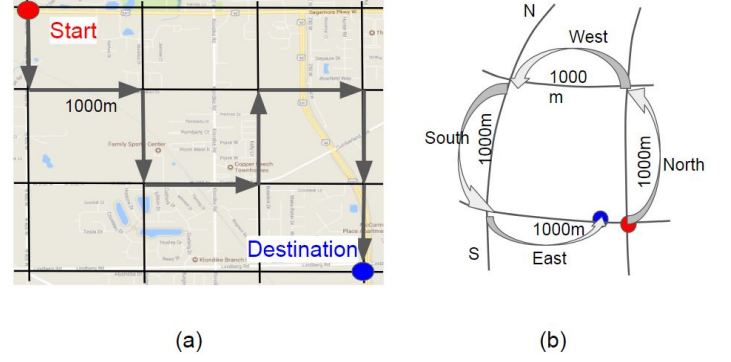

(a)   (b)

Figure 5: (a) Grid map of environment which the agent navigate on. The length and the direction of each arrow represents a certain length of displacement for a step and action taken by the agent. (b) In the real world, the agent will not be back to the same position if it take a round sequence of steps. This is caused by the sphere geometry.

Interact with Google Map API We first specify the start position (could be a place name, address or geocode) and the destination position and input to the Geocoding API to retrieve geocodes of these two position. We then use the start geocode and destination geocode as the two opposite corner to construct the rectangle boundary of the grid map where our agent can only navigate on. The grid map and the symbols are shown in figure 5(a). Strictly speaking, each grid in the grid map is not a rectangle. This phenomenon is caused by the sphere geometry and our restriction on the length of the stride which is demonstrated in figure 5(b).

There are four directions choices (north, east, south and west) for the agent. Each arrow represents the agent navigating from the current position (s) to the next position (s') with 1000 meters of displacement on the grid map. However, the real navigating distance of the agent will be larger or equal to 1000m depends on the route provided by the Directions API. For example, in figure 6, assume that the agent is at current position denoted by A and heading south to the next position denoted by B. Apparently, the route provided by the Directions API is the highway 395 and the distance is longer than 1000m. We can get the navigating instruction list with the form: { geocode of A, duration from A to 1, distance from A to 1, geocode of 1 }, { geocode of 1, duration from 1 to 2, distance from 1 to 2, geocode of 2 } .... where A, 1, 2 ,3 , B are shown in figure 6 after we input the geocode of A and the geocode of B into the Directions API. The number of instruction is based on the Directions API and there are four instructions in our case of figure 6 from A to B. We use each of the geocode in the navigating instruction list to get the height of each position from the Elevation API and compute the elevation within each instruction such as the elevation between A(position 1 in figure 7) and 1(position 2 in figure 7), the elevation between 1 and 2, the elevation between 2 and 3, and the elevation between 3 and B. We omit the height in the middle of the road between two position provided within each instruction to simplify the complexity
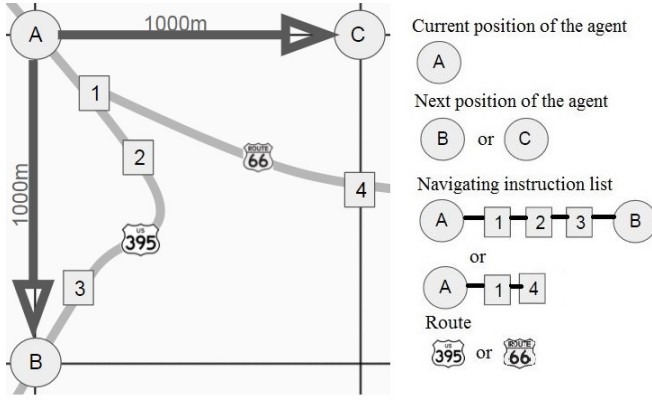
Figure 6: The graph only shows a part of the grid map. The agent can only navigate from one point to the other nearby point without crossing the boundary. We will analyze the road information such as road 395 if the agent travel from the current position A to the next position B.
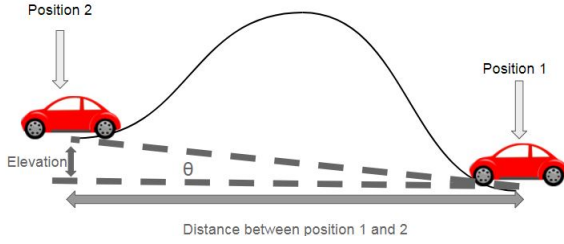

Figure 7: We only compute the elevation between position 1 and 2 and neglect the elevation on the way

shown in figure 7. However, the case for the agent to travel from A to C will depends on the information that Directions API returns. If C is in a lake or somewhere unreachable, the Directions API will return false and the agent will regards the point C as a block. Otherwise, position C is considered reachable and the navigation instruction list returned by the Directions API is route 66, then we will analyze the elevation between position A, 1 and 4 even though the position 4 is not the identical position as C. We allow the route returned by the Directions API is located outside the grid map boundary while the agent should always navigate on the point within the boundary of the grid map.

Energy Consuming   We evaluate each action based on the energy required to travel with 1000m displacement on the grid map (real distance is more than 1000m). To compute the energy required between position A and position 1 in figure 6, for example, we use the duration and distance between position A and position 1 to calculate the average velocity V. Combine V with the elevation, we can get the angle θ of the road and consider the height of the road as linear increasing or decreasing shown in Figure 7. Because we don't take regenerative braking into account in our experiment, we treat the downhill road flat. In figure 8, we demonstrate how do we calculate the power required(Garcia-Valle and Lopes 2012) for a car with mass M (kg) to travel on the road with
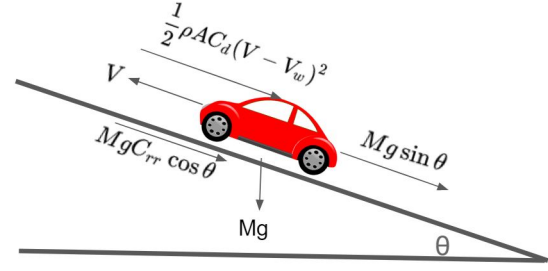

Figure 8: Free-body diagram of forces applied on a vehicle

angle θ (degree) under velocity V (m/s) is as follows

$$P = f_m M\alpha V + MgC_{rr}V\cos(\theta) + \frac{1}{2}\rho AC_d(V - V_W)^2 + MgV\sin(\theta)$$

where P (W) is power, f_m is the mass factor, M (kg) is the overall mass, g (m/s^2) is the acceleration of gravity, C_{rr} is the coefficient of rolling resistance between tires and road surface, ρ is the air density (kg/m^3), A (m^2) is the vehicle frontal area, C_d is the aerodynamic drag coefficient and V_W is the wind speed. We don't consider the early stage of acceleration, so α is zero. The parameters are listed in Table 1. Energy consumption should be computed by multiplying the power P by the duration.

Table 1: Parameters for power calculation

f_m	1.05
α	0 m/s^2
Mass	2000 kg
C_{rr}	0.02
ρ	1.225 kg/m^3
A	2 m^2
C_d	0.5
V_W	0 m/s

Reward Arrangement   The fundamental concept of defining the reward is based on the energy consumption in one stride from the current position to the next position, for example, from A to B shown in the figure 6. The energy is calculated by the method provided previous section. We then divide the energy by 10000 and times -1. In order to minimize the number of total steps during training, we add -0.1 to each transition if the next position is reachable. In other words, the reward r for taking any reachable step will be r = -0.1 - (energy consumption / 10000). If the next position is unreachable such as a lake or a river, r = -1 and the agent stay at the same current position and take the other action. If the distance of the next position and the destination position is less than the length of the predefined displacement (1000m in the figure 6), the reward r for taking this action will become r = +1 - (energy consumption from the current position to the next position / 10000) - (energy consumption from the next position to the destination position / 10000). Noticed here +1 appears in the reward because of the success of this action which lead the agent to the destination.
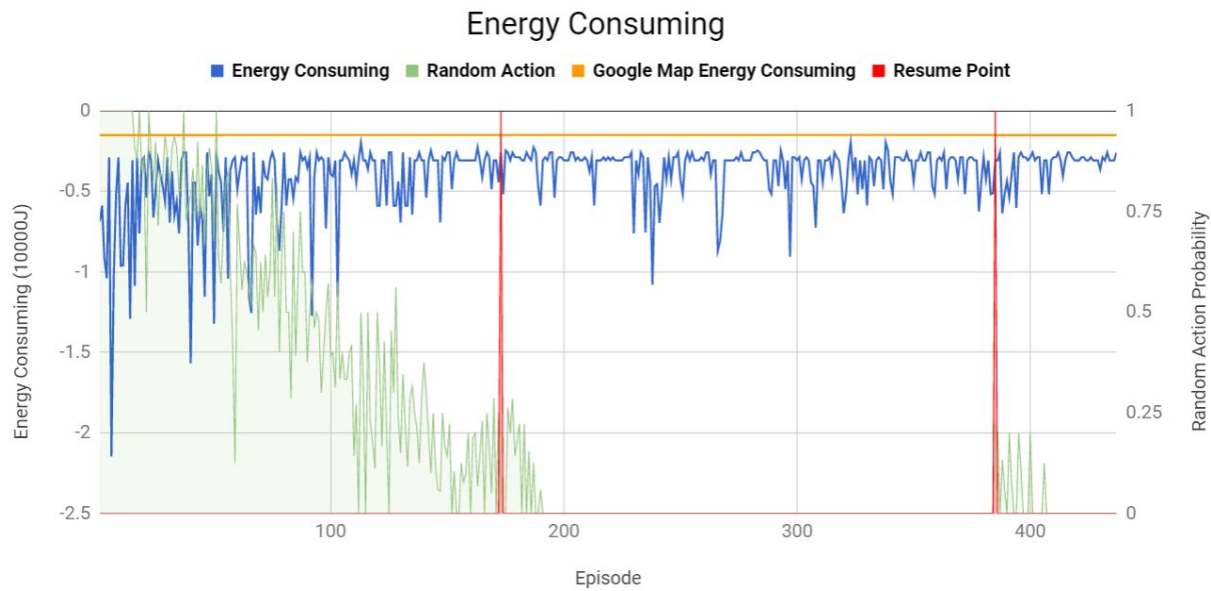
Energy Consuming

Figure 9: The energy consuming of the learning agent. The blue line is the energy consumed and the orange line is the energy required by taking the route which Google map recommend. The green line represents the probability of taking action randomly and the red line means we resume the training from previous training model

Electrical Vehicle Battery   In the experiment, the battery performance will not affect the training process. The battery is able to carry totally 50000Wh of energy which is a standard offering by electrical vehicle manufacture, Tesla. In electrochemistry, it is recommended to use the the state of charge (SOC) from $90\% \sim 20\%$ of a battery to improve it's life which we also implement in our case. The SOC is calculated by the ratio between the current energy and the total energy. We will not take the battery degradation into the experiment. Further work can take the real factor on battery performance into account as part of the training process. For this experiment, we only demonstrate how much energy consumed and how many times the battery need to be charged in an ideal condition.

Results and Discussion

The agent is being trained from the start position (geocode:40.4682572, -86.9803475) and the destination (geocode:40.445283, -86.948429) with stride length 1000m. Steps more than 36 steps within an episode will be regarded as failed. Noticed that during the training, Google map api often blocked our server and we are forced to end the training process and resume the model from the interrupted episode (red line). This problem will lead to the empty replay buffer where we choose sample uniformly at random to compute the loss and updtate the weights during learning. As a result, we will need to resume the model and start choosing action randomly to refill the replay buffer and gradually decrease the portion of random action.

The blue line in figure 9 shows the energy consumed by the agent. The agent is able to find out a way to minimize the energy consumption after 100 episodes. The oscillation in the first 100 episodes is caused by highly random ac-

tion (green line) and inaccurate Q value provided by the Q-network. While the loss of inaccurate Q value is minimized, the oscillation is mitigated and the energy consumption become less and stable. The minimum energy that the agent can achieve without random action is 1951(J) and the corresponding time is 946 seconds where the energy and time of the route provided by Google map are 1489(J) and 315 seconds.

Conclusion

Double-DQN is able to learn from the reward and experiences. Decrease the length of the stride to 20m or 50m can increase the accuracy and probably the energy reduction but still required more experiments to confirm. Noticed that the more accuracy you can obtain, the more computation resources and time you will need. Source code: https://github.com/Dungyichao

References

Garcia-Valle, R., and Lopes, J. A. P. 2012. *Electric vehicle integration into modern power networks*. Springer Science & Business Media.

Tamar, A.; Wu, Y.; Thomas, G.; Levine, S.; and Abbeel, P. 2016. Value iteration networks. In *Advances in Neural Information Processing Systems*, 2146–2154.

Van Hasselt, H.; Guez, A.; and Silver, D. 2016. Deep reinforcement learning with double q-learning. In *AAAI*, 2094–2100.

Wang, X., and Gupta, A. 2016. Generative image modeling using style and structure adversarial networks. In *European Conference on Computer Vision*, 318–335. Springer.